

Using C from Python

DesertPy - Phoenix Python Meetup Group

Dave Kaspar

March 27, 2019



Outline

- 1 Example: Truncated search of numeric array
- 2 Options for calling native code
- 3 Overview of `ctypes`
- 4 Truncated search in `C`, `ctypes` wrapper
- 5 More involved usage

Consider the following task

- Given a vector $\mathbf{x} = (x_0, \dots, x_{L-1})$ of length L
- Given a real value v
- Find $i = \min\{j : x_j < v\}$ or determine that no such i exists

It's trivial to implement this in Python:

```
def firstless(x, v):  
    L = len(x)  
    i = 0  
    while i < L and x[i] >= v:  
        i += 1  
    return i
```

Here we use the out-of-bounds index $i = L$ to indicate no such i exists; we could also return `None`

What's wrong with this?

```
def firstless(x, v):  
    L = len(x)  
    i = 0  
    while i < L and x[i] >= v:  
        i += 1  
    return i
```

Algorithm is entirely satisfactory, but:

- Imagine this tight loop sits inside another loop, and is the hottest part of the code (not hypothetical)
- VM overhead and support of dynamic language features will dominate execution time

Built-ins less than ideal

Alternative implementation using numpy:

```
import numpy as np

def firstless_overworked(x, v):
    x = np.asarray(x)
    nz = np.nonzero(x < v)[0]
    return (nz[0] if len(nz) else len(x))
```

- The good: array operations call out to native code
- $x < v$ is a Boolean array (same dimensions as x)
- `np.nonzero($x < v$)` returns a 1-tuple holding an array of indices where $x < v$ is **True**
- Algorithm is now silly, iterates through two arrays of full length `len(x)`

Truncated search example

- Was deliberately simple (but not contrived)
- Illustrates a case where writing and calling native code offers substantially improved efficiency (if this is needed)

Outline

- 1 Example: Truncated search of numeric array
- 2 Options for calling native code**
- 3 Overview of `ctypes`
- 4 Truncated search in `C`, `ctypes` wrapper
- 5 More involved usage

Option: Python API

We could use Python's C API to write an extension module

- **GOOD**: Layer between C and Python is thin from the perspectives of:
 - Python users (can supply docstrings, function signatures, ...)
 - Execution time
- **GOOD**: Compile, install, distribute using distutils/setuptools
- **GOOD**: Mature option, documentation is pretty good
- **BAD**: Layer between C and Python is not thin in the sense of programmer time:
 - `PyArg_ParseTuple()`
 - Manual reference counting
 - Write algorithm in generic C, then write C wrappers for (each?) dynamic language of interest

Verdict: Suitable in certain situations but not in generality

Option: SWIG

Wrapper generator for C, C++ libraries producing bindings for *many* dynamic languages

- **GOOD**: Target several languages
- **BAD**: One size does not fit all
- **BAD**: Need SWIG *.i interface files

Verdict: May be appropriate for a substantial library, with a simple object model, when use from several dynamic languages is anticipated

Option: Boost.Python

This C++ library provides macros and functions for exposing C++ functions, classes, namespaces in Python extension modules

- **GOOD**: Specializing on just the pair (C++, Python) enables better matching of object models than SWIG could hope to achieve
- **BAD**: Limited to just C++ and Python
- **BAD**: The boilerplate to be written must be written in C++

Verdict: If you're using C++, this seems like the best bet

Option: Cython

```
# cyfirstless.pyx
cimport cython

@cython.boundscheck(False)
@cython.wraparound(False)
def firstless(double [:] x, double v):
    cdef size_t i = 0
    cdef size_t L = len(x)
    while i < L and x[i] >= v:
        i += 1
    return i
```

```
# setup.py
from setuptools import setup
from Cython.Build import cythonize

setup(ext_modules = cythonize("cyfirstless.pyx"))
```

Option: Cython

Cython transpiles to C, then compiles C to Python extension in shared object *.`(so|dylib|dll)`

- **GOOD**: Start writing nearly standard Python, optionally add type information
- **GOOD**: Interacts nicely with `numpy`
- **BAD**: Limited support for parallelism (essentially OpenMP for loops)
- **BAD**: Not portable to R, Matlab (if you care about these)

Verdict: An excellent option in many circumstances, and probably best for the truncated search example

Option: ctypes

The ctypes module from the standard library provides support for C data types (including aggregates) and calling functions in shared objects

- **GOOD**: No additional dependencies
- **BAD**: There is some boilerplate
- **GOOD**: There's not much of it if you're only using a few functions
- **GOOD**: You write the boilerplate in Python
- **BAD**: (Essentially) limited to C, Python

Verdict: An excellent option for one-off library calls or when you want to exercise a lot of control over the Python interface

Outline

- 1 Example: Truncated search of numeric array
- 2 Options for calling native code
- 3 Overview of ctypes**
- 4 Truncated search in C, ctypes wrapper
- 5 More involved usage

- Classes for fundamental C data types, examples:
 - `c_int`
 - `c_double`
 - `c_char_p` (not quite fundamental, but gets special treatment)
- Classes for aggregates, derived types:
 - Structure – subclass to define your `structs`
 - Arrays – `c_int * 16` is the *type* of an array of 16 C integers
 - Pointers – `POINTER(T)` is the *type* of a pointer to type T
- Each object constructed from one of these classes is backed by a memory buffer holding the equivalent C data, can pass to C functions `byref()`
- The values of fundamental types are accessible, settable via the `.value` `property`
- Each field of a structure is also `property`, gettable and settable

ctypes libraries and their functions

- Functions are available as attributes of libraries
- We specify return and argument types
- Then ctypes handles type conversions for us

```
from ctypes.util import find_library
from ctypes import * # for brevity here

libm = CDLL(find_library("m"))
libm.sqrt.restype = c_double
libm.sqrt.argtypes = (c_double,)

sqrt_5 = libm.sqrt(5)
print(type(sqrt_5), sqrt_5)
# prints <class 'float'> 2.23606797749979
```


ctypes callbacks

Most of the time we are interested in calling C from Python, but sometimes C functions take a function pointer as an argument.

```
from ctypes import * # for brevity here
```

```
# type of C function taking char * argument, returning size_t  
# is CFUNCTYPE(c_size_t, c_char_p), can be used as decorator
```

```
@CFUNCTYPE(c_size_t, c_char_p)  
def strlen(b):  
    return len(b.decode())
```

We can pass this strlen to a C function taking an argument

```
size_t (*func)(char *).
```

Outline

- 1 Example: Truncated search of numeric array
- 2 Options for calling native code
- 3 Overview of ctypes
- 4 Truncated search in C, ctypes wrapper**
- 5 More involved usage

Truncated search in C (obvious implementation)

```
/* cfirstless.c */
#include <stddef.h>
#include <stdint.h>
#include "cfirstless.h"

size_t
cfirstless(const double *x, size_t len, double v)
/* Given array of length len at address x, return the first
 * index i for which x[i] < v, or return len if no such
 * index exists. */
{
    size_t i = 0;
    while (i < len && x[i] >= v)
        i += 1;
    return i;
}
```

Truncated search in C (blocked implementation)

```
typedef double v8df __attribute__((vector_size (64)));
typedef int64_t v8di __attribute__((vector_size (64)));

size_t
cfirstless_blocked(const double *x, size_t len, double v)
/* Given array of length len at address x, return the first
 * index i for which x[i] < v, or return len if no such
 * index exists. Process in blocks of 8 to facilitate SIMD
 * optimizations. */
{
    size_t i, j;
    size_t len_first = ((uintptr_t) x) & 63;
    if (len < len_first)
        len_first = len;
    for (i = 0; i < len_first; i++)
        if (x[i] < v)
            return i;
```

Truncated search in C (blocked implementation)

```
v8df *xv;
v8df vv = {v, v, v, v, v, v, v, v};
v8di cmp;
while (i + 8 < len)
{
    xv = __builtin_assume_aligned(&x[i], 64);
    cmp = (*xv) < vv;
    for (j = 0; j < 8; j++)
        if (cmp[j])
            return i + j;
    i += 8;
}

while (i < len && x[i] >= v)
    i += 1;
return i;
}
```

Compilation to shared library

```
gcc -O3 -march=native -fpic -Wall -shared \  
-o libcfirstless.so cfirstless.c
```

Python wrapper firstless

```
# firstless.py
import ctypes as C
from warnings import warn
import numpy as np
import cyfirstless

try:
    lcfl = C.CDLL("./libcfirstless.so")
    for f in (lcfl.cfirstless, lcfl.cfirstless_blocked):
        f.restype = C.c_uint
        f.argtypes = (C.POINTER(C.c_double),
                      C.c_uint, C.c_double)
        # print("Success loading libcfirstless.so")
except:
    warn("Failed to load libcfirstless.so")
    lcfl = None
```

Python wrapper firstless

```
def firstless(x, v, method = "C"):
    x = np.ascontiguousarray(x, dtype = float)
    method = method.upper()
    if method in ("C", "B") and lcf1:
        px = x.ctypes.data_as(C.POINTER(C.c_double))
        if method == "C":
            return lcf1.cfistless(px, len(x), v)
        else:
            return lcf1.cfistless_blocked(px, len(x), v)
    elif method == "CY":
        return cyfirstless.firstless(x, v)
    elif method == "N":
        nz = np.nonzero(x < v)[0]
        return (nz[0] if len(nz) else len(x))
```


Python wrapper firstless

```
else:
    if method != "P":
        warn("firstless falling back to pure python")
    i = 0
    while i < len(x) and x[i] >= v:
        i += 1
    return i
```

```
def time():
    from timeit import repeat
    stmt = "firstless(x, v, method = '{}')"
    setup = """import numpy as np
x = np.arange(10 ** 6, -1, -1)
v = x[int(0.25 * 10 ** 6)]"""
```

Python wrapper firstless

```
desc = dict(C = "C (obvious implementation)",
            B = "C (blocked implementation)",
            CY = "Cython",
            N = "Numpy (built-ins, overworking)",
            P = "Python (purely)")
for m in desc:
    t = repeat(stmt.format(m), setup,
               number = 1000, globals = globals())
    print("Method {} 1000 times: {}".format(desc[m], t))
```

Time results (old machine)

```
In [3]: firstless.time()
```

Method C (obvious implementation) 1000 times:

```
[4.564280847000191, 4.561867651005741, 4.455790193998837]
```

Method C (blocked implementation) 1000 times:

```
[4.508035304999794, 4.651949233004416, 4.607617080000637]
```

Method Cython 1000 times:

```
[4.4173163150044275, 4.484859807002067, 4.680379589000095]
```

Method Numpy (built-ins, overworking) 1000 times:

```
[8.515583603999403, 8.81348213399906, 9.226202923993696]
```

Method Python (purely) 1000 times:

```
[139.9592333890032, 139.45108247499593, 139.2754626599999]
```

Time results (newer machine)

```
In [3]: firstless.time()
```

Method C (obvious implementation) 1000 times:

```
[1.1181618470000103, 1.1126245070190635, 1.1092812279821374]
```

Method C (blocked implementation) 1000 times:

```
[1.1583812900062185, 1.1572002400062047, 1.1530515620252118]
```

Method Cython 1000 times:

```
[1.104163747979328, 1.1027535720204469, 1.1014864780008793]
```

Method Numpy (built-ins, overworking) 1000 times:

```
[1.980167209985666, 1.981402239005547, 1.9776558729936369]
```

Method Python (purely) 1000 times:

```
[45.68740793998586, 43.17407809701399, 43.217961503978586]
```

Outline

- 1 Example: Truncated search of numeric array
- 2 Options for calling native code
- 3 Overview of `ctypes`
- 4 Truncated search in `C`, `ctypes` wrapper
- 5 More involved usage**

The reason I needed ctypes

I've been working on a project which involves both:

- Embarrassingly parallel operations
- Shared structure

More specifically:

- Audio and symbolic segmentation using HMM ideas and dynamic programming
- Partial segmentations form a tree, good branches extended, stale branches eventually discarded

Want:

- Work in one address space
- Avoid GIL
- Drive things interactively from a REPL

A C library exposed to Python via ctypes achieves these objectives

ctypes issues arising in this work

- Who owns the memory?
 - We might care about use outside Python, so the C library does its own allocation
 - This means a few Python classes have `__del__()` methods that call `*_free()` functions in the C library
- Original object return – ctypes doesn't do this
 - When you extract the value of a ctypes object, or dereference a pointer, ctypes constructs a new object for the result
 - Usually not a problem if you're just reading results
- Flexible array members
 - Multiple levels of segmentation, C code aims to be generic
 - Generic `structs` have a buffer to hold specific data, the buffer is a flexible array member
 - ctypes retrieves header of `struct`, from which we can get buffer size
 - Can resize buffer backing the retrieved ctypes Structure, then `memcpy` the rest of the data

Conclusion

- Need sometimes arises to write/call native code
- Several options for connecting this to Python
- With `ctypes` you define the interface in Python, suitable both for simple situations and cases where you want a lot of control
- Happy to answer any questions
- Thank you for your attention!